

Systems Engineering of LLM Workstation

Kirill Nevzorov under supervision
of Prof. Christopher Buckley
University of Sussex
School of Engineering and Informatics

February 21, 2026

Abstract

We wish to research the process of creating and use-case testing various LLM prompts, such as prompting the user for additional details relevant to their project, assisting them in resummarising and elaborating on the information stored within a workspace, and generating intelligence from a collection of notes.

Contents

I. Introduction	2	4.2.2 Error Handling and Recovery . . .	3
1.1 Background and Motivation	2	4.2.3 Logging and Debugging	3
1.2 Project Scope and Objectives	2	4.3 Security and Access Control	3
1.3 Report Structure	2	V. Research Findings and Analysis	4
II. Foundations of LLM Systems	2	5.1 LLM Router and Load Balancer Evaluation	4
2.1 Understanding Large Language Models .	2	5.1.1 llama.cpp Ecosystem Exploration .	4
2.1.1 Model Architecture Overview . . .	2	5.1.2 Ollama Architecture Insights . . .	4
2.1.2 Inference vs. Training Workloads .	2	5.1.3 Performance Comparisons	4
2.2 LLM Workload Characteristics	2	5.2 System Design Tradeoffs	4
2.2.1 Latency Requirements	2	5.2.1 Cost vs. Performance	4
2.2.2 Throughput Considerations	2	5.2.2 Complexity vs. Maintainability . .	4
2.2.3 Resource Utilization Patterns . . .	2	5.2.3 Flexibility vs. Stability	4
2.3 System Integration Challenges	2	5.3 Lessons Learned from Exploration	4
III. Architectural Patterns and Design	3	VI. Conclusion and Future Work	4
3.1 LLM Inference Serving Architectures . . .	3	6.1 Summary of Key Findings	4
3.1.1 Single Model Serving	3	6.2 Recommendations for Implementation . .	4
3.1.2 Multi-Model Ensembles	3	6.3 Areas for Further Research	4
3.1.3 Model Routing and Load Balancing	3	Appendices	5
3.2 Backend Selection Strategies	3	A. Code Examples and Snippets	5
3.2.1 Hardware Considerations	3	B. System Diagrams and Architecture Charts	5
3.2.2 Software Stack Evaluation	3	C. Research Resources and References	5
3.2.3 Cost-Benefit Analysis	3		
3.3 Query Management and Optimization . .	3		
3.3.1 Request Queuing and Prioritization	3		
3.3.2 Caching Strategies	3		
3.3.3 Batch Processing Approaches . . .	3		
3.4 Scalability and Performance Patterns . .	3		
IV. Practical Implementation Considerations	3		
4.1 Infrastructure and Deployment	3		
4.1.1 Containerization and Orchestration	3		
4.1.2 Cloud vs. On-Premise Considerations	3		
4.1.3 Resource Provisioning	3		
4.2 Monitoring and Observability	3		
4.2.1 Performance Metrics	3		

I. Introduction

1.1 Background and Motivation

In order to build useful software. . .

1.2 Project Scope and Objectives

We wish to interact with Large Language Models efficiently. . .

1.3 Report Structure

This report is organised into six main chapters, followed by appendices:

Chapter I: Introduction provides the background and motivation for this project, defines the scope and objectives, and outlines the structure of the report.

Chapter II: Foundations of LLM Systems establishes the theoretical groundwork by examining large language model architectures, distinguishing between inference and training workloads, characterising LLM workload requirements in terms of latency, throughput, and resource utilisation, and identifying key system integration challenges.

Chapter III: Architectural Patterns and Design explores the design space for LLM serving systems, covering inference serving architectures (single model, multi-model ensembles, and model routing), backend selection strategies (hardware, software stack, and cost-benefit considerations), query management and optimisation techniques (queuing, caching, and batch processing), and scalability patterns.

Chapter IV: Practical Implementation Considerations addresses the engineering aspects of deploying LLM systems, including infrastructure and deployment (containerisation, cloud vs. on-premise, resource provisioning), monitoring and observability (metrics, error handling, logging), and security and access control.

Chapter V: Research Findings and Analysis presents the results of our exploration, including evaluations of LLM routers and load balancers (with specific focus on the llama.cpp ecosystem and Ollama architecture), system design tradeoffs (cost vs. performance, complexity vs. maintainability, flexibility vs. stability), and lessons learned.

Chapter VI: Conclusion and Future Work summarises the key findings, provides recommendations for implementation, and identifies areas for further research.

Appendices include code examples and snippets, system diagrams and architecture charts, and research resources and references.

II. Foundations of LLM Systems

2.1 Understanding Large Language Models

2.1.1 Model Architecture Overview

2.1.2 Inference vs. Training Workloads

2.2 LLM Workload Characteristics

2.2.1 Latency Requirements

2.2.2 Throughput Considerations

2.2.3 Resource Utilization Patterns

2.3 System Integration Challenges

III. Architectural Patterns and Design

3.1 LLM Inference Serving Architectures

3.1.1 Single Model Serving

3.1.2 Multi-Model Ensembles

3.1.3 Model Routing and Load Balancing

3.2 Backend Selection Strategies

3.2.1 Hardware Considerations

3.2.2 Software Stack Evaluation

3.2.3 Cost-Benefit Analysis

3.3 Query Management and Optimization

3.3.1 Request Queuing and Prioritization

3.3.2 Caching Strategies

3.3.3 Batch Processing Approaches

3.4 Scalability and Performance Patterns

IV. Practical Implementation Considerations

4.1 Infrastructure and Deployment

4.1.1 Containerization and Orchestration

4.1.2 Cloud vs. On-Premise Considerations

4.1.3 Resource Provisioning

4.2 Monitoring and Observability

4.2.1 Performance Metrics

4.2.2 Error Handling and Recovery

4.2.3 Logging and Debugging

4.3 Security and Access Control

V. Research Findings and Analysis VI. Conclusion and Future Work

5.1 LLM Router and Load Balancer Evaluation

5.1.1 llama.cpp Ecosystem Exploration

5.1.2 Ollama Architecture Insights

5.1.3 Performance Comparisons

5.2 System Design Tradeoffs

5.2.1 Cost vs. Performance

5.2.2 Complexity vs. Maintainability

5.2.3 Flexibility vs. Stability

5.3 Lessons Learned from Exploration

6.1 Summary of Key Findings

6.2 Recommendations for Implementation

6.3 Areas for Further Research

Appendices

A. Code Examples and Snippets

B. System Diagrams and Architecture
Charts

C. Research Resources and References